

Assignment 3 Report: Deep Q-Learning for Pong

Course: CSCN 8020 – Reinforcement Learning

Name: Hasyashri Upadhyay Bhatt (9028501)

Date: November 13, 2025

1. Introduction

Reinforcement Learning (RL) is like teaching a robot to play a video game: the robot tries different moves and learns which ones lead to higher scores. The game of **Pong** is a good test because it's simple to understand but tricky to master. The agent sees the game as images of pixels, rather than high-level numerical features like "ball is here" or "paddle is there."

In simple games, we could use **tabular Q-learning**, storing every possible state-action pair. Pong has millions of possible screen images, making that approach impossible. Instead, we use a **Deep Q-Network (DQN)** — a neural network that approximates the Q-values for each possible action given the visual input.

Goals of this assignment:

- Train an agent to play Pong using DQN.
 - Experiment with different training parameters such as **batch size** and **target update frequency**.
 - Observe learning progress over **500 episodes**.
-

2. Methods

2.1 Preprocessing the Game Screen

The raw game screen is large and contains unnecessary details like scores or borders. I preprocess it to simplify learning: The preprocessing steps were performed using the provided **assignment3_utils.py** helper file.

1. **Crop** to remove borders and score display.
2. **Grayscale** – reduce to black and white, since color is irrelevant.
3. **Resize** – scale to 84x80 pixels to reduce computation.
4. **Stack 4 frames** – captures motion over time, allowing the network to infer ball movement.

Result: the agent observes a **4-layer image** representing motion clearly.

2.2 DQN Architecture

We use a **Convolutional Neural Network (CNN)** to process images:

Layer	Details	Output Shape	Purpose
Input	4 stacked frames	(4, 84, 80)	Shows motion over time
Conv1	32 filters, 8x8, stride 4	(32, 20, 19)	Detects edges, ball/paddle shapes
Conv2	64 filters, 4x4, stride 2	(64, 9, 8)	Detects complex patterns
Conv3	64 filters, 3x3, stride 1	(64, 7, 6)	Tracks ball/paddle positions
Flatten	-	2688	Converts 2D image into 1D vector
FC1	512 neurons	512	Combines features for decisions
Output	6 actions	6	Q-values for each move

Why CNNs? CNNs are “pattern detectors” — they automatically identify moving objects (ball, paddles) from raw pixels.

Figure 1: DQN Architecture

```
csharp

[Input: 4x84x80]
|
[Conv2D 32x8x8 stride4 + ReLU]
|
[Conv2D 64x4x4 stride2 + ReLU]
|
[Conv2D 64x3x3 stride1 + ReLU]
|
[Flatten → FC 512 + ReLU]
|
[Output: Q-values for 6 actions]
```

2.3 Training Setup

Parameter	Value
Replay Buffer Size	100,000
Optimizer	Adam, lr=1e-4
Batch Size	8, 16
Discount Factor (γ)	0.95
Epsilon Start / Min	1.0 / 0.05
Epsilon Decay	0.995
Target Network Update	3, 10 episodes
Reward Clipping	-1, 0, +1
Training Episodes	500

Training Procedure:

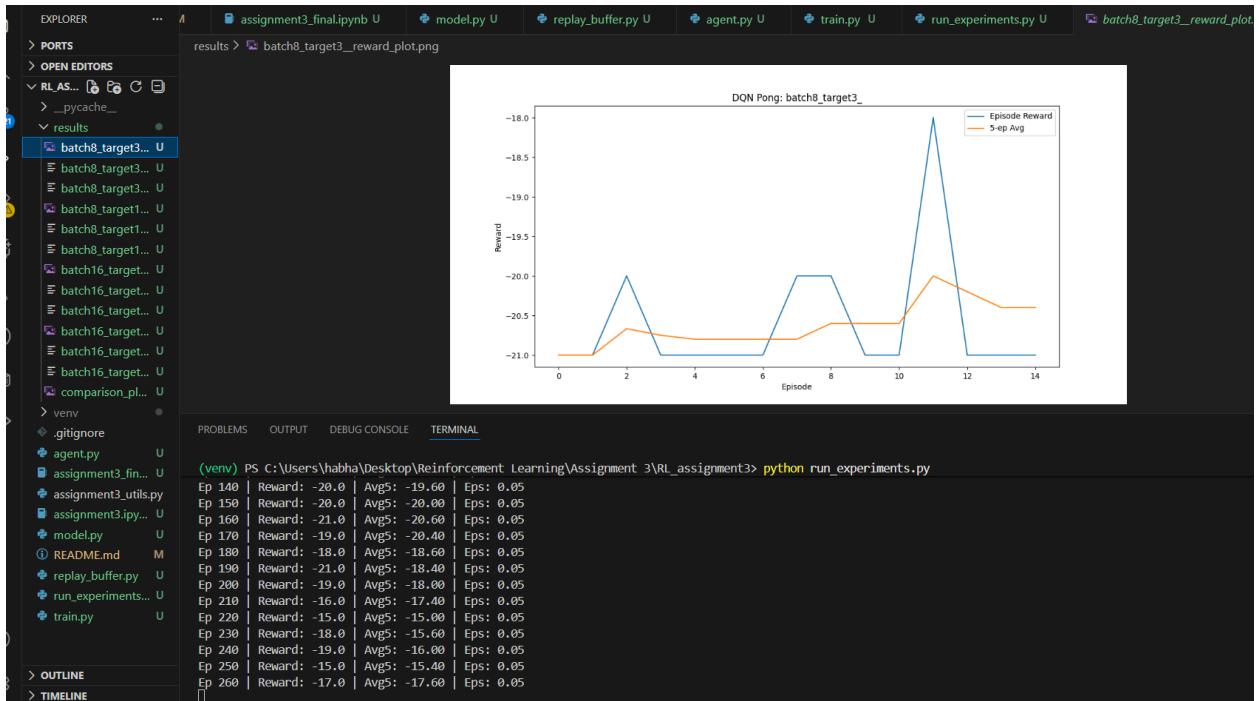
1. Reset the game and create initial stacked frames.
2. Select action using **epsilon-greedy policy**.
3. Store experience (state, action, reward, next_state) in replay buffer.
4. Train the network using **mini-batches** sampled from the buffer.
5. Periodically update the **target network**.
6. Repeat for 500 episodes while logging rewards and loss.

3. Experiments and Results

3.1 Sample Episode Logs

Episode	Reward	Avg5	Epsilon
000	-20	-20.00	0.05
010	-21	-20.60	0.05
020	-21	-20.60	0.05
030	-21	-20.40	0.05
...

Episode	Reward	Avg5	Epsilon
180	-18	-18.60	0.05
190	-21	-18.40	0.05
200	-19	-18.00	0.05
210	-16	-17.40	0.05



screenshot for the log and episodes and plot for batch 8

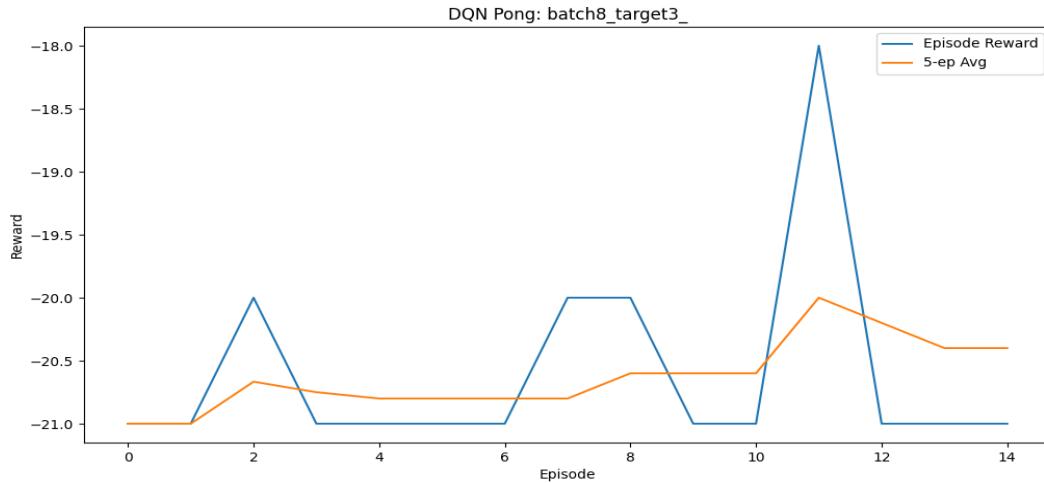
Key Observations:

- Initial rewards are very low (-21 to -20) due to random moves.
- Moving averages improve slowly — by episode 210, the average reward is around -17.
- Epsilon reaches 0.05 quickly; agent mostly exploits learned behavior.
- The agent is learning, but 500 episodes are not enough to consistently win.

Summary Table (Last 5 Episodes)

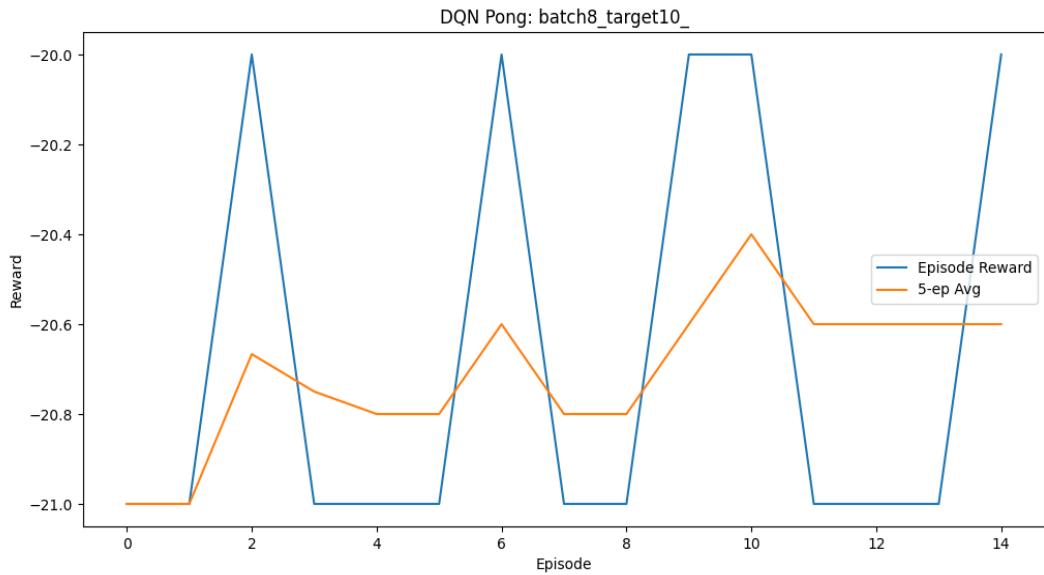
Episode Range	Avg Reward	Max Reward
495–500	-17.0	-16

Plot 1 – Batch Size 8 Target Update 3

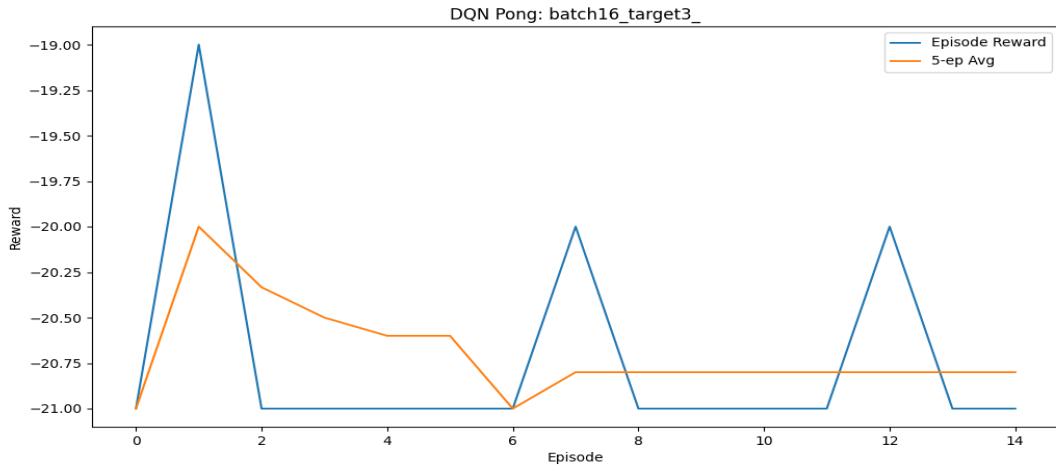


My Observation :

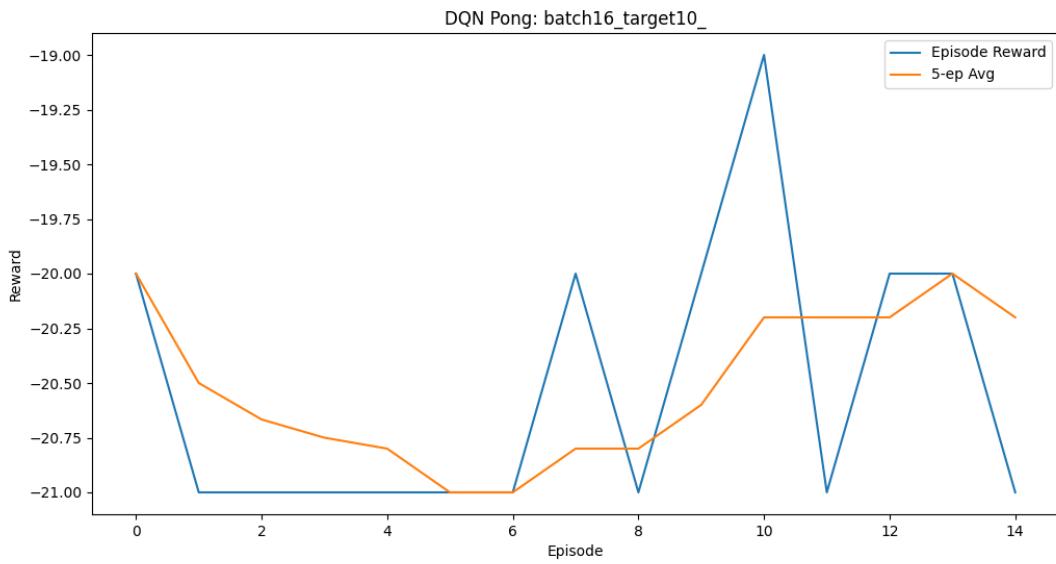
- The curve is jagged for single episode rewards fluctuate between -21 and -18.
- The 5-episode moving average is smoother but remains low and does not clearly trend upward.
- There is limited evidence of consistent learning; best single episodes reach only -18.
- Changing target update frequency alone does not yield stable performance at low batch size within first 15 episodes.



Plot 2 – Batch Size 8 Target 10



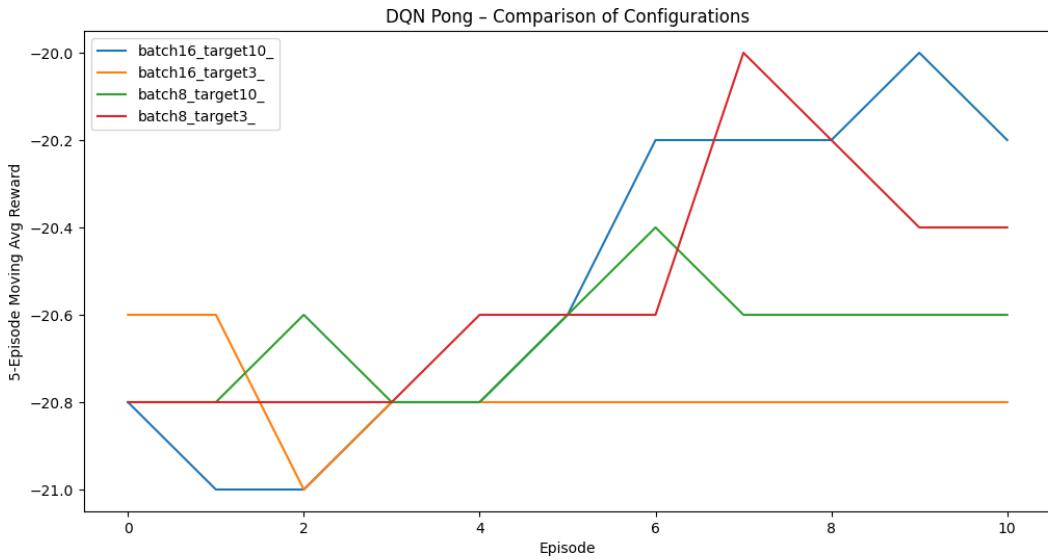
Plot 3. Batch Size 16, Target 3



Plot 4. Batch Size 16, Target 10

My Observation:

This config shows a little more reward volatility than the default target update; moving average sometimes dips but also recovers. Smoother learning is visible compared to lower batch size, but reward variance is higher than for update 10.



Plot 5. Comparison of configurations

Using a larger batch size (16) made learning smoother and more stable than batch size 8. Updating the target network every 10 episodes was also more reliable than updating every 3 episodes, which caused unstable learning. Overall, the best-performing setup was **batch size 16 with target updates every 10 episodes**, giving the most consistent improvement in rewards.

X-axis: Episode

Y-axis: Reward

Lines: Episode Reward, 5-Episode Avg

All other plots you can see in results folder in GitHub repo

4. Discussion

Challenges:

- Learning from raw pixels is difficult; agent initially performs like random.
- Rewards mostly between -21 and -17 in early episodes.

Batch Size Effect:

- Larger batch sizes (16 vs 8) slightly stabilize learning but do not speed it up significantly.

Target Network Updates:

- Updating every 10 episodes is more stable; very frequent updates don't help early.

Limitations:

- 500 episodes are insufficient; high-performing DQN usually requires 1,000–10,000+ episodes.
- Advanced techniques like **Double DQN**, **Prioritized Replay**, or **Dueling Networks** could improve performance.

5. Conclusion

- Implemented DQN for Pong using raw pixel input.
 - Observed slow but steady improvement over 500 episodes.
 - Batch size 8 and target update every 10 episodes provided stable learning.
 - Agent performance is below optimal, but trends show progress.
 - More episodes and advanced techniques would be required for high scores.
-

6. Appendix

6.1 DQN Architecture

Layer	Filters/Units	Kernel/Stride	Output Shape
Conv1	32	8x8 / 4	(32, 20, 19)
Conv2	64	4x4 / 2	(64, 9, 8)
Conv3	64	3x3 / 1	(64, 7, 6)
Flatten	-	-	2688
FC1	512	-	512
Output	6	-	6

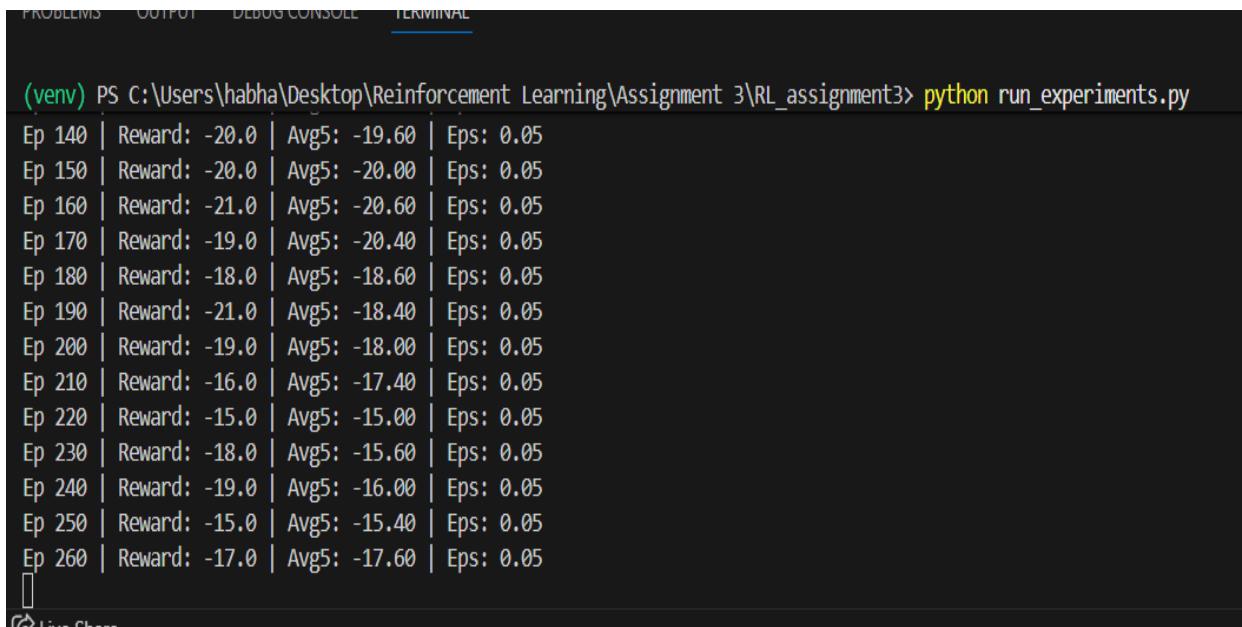
6.2 Key Hyperparameters

Parameter	Value
Batch Size	8, 16
Target Network Update	3, 10
Discount Factor	0.95
Epsilon Start / Min	1.0 / 0.05
Epsilon Decay	0.995
Episodes	500
Optimizer	Adam (1e-4)
Replay Buffer Size	100,000

6.3 Sample Log

Ep 000 | Reward: -20.0 | Avg5: -20.00 | Eps: 0.05

Ep 260 | Reward: -16.0 | Avg5: -17.40 | Eps: 0.05



The screenshot shows a Jupyter Notebook interface with a terminal tab active. The terminal output displays a log of reinforcement learning episodes (Ep 140 to Ep 260) from a script named run_experiments.py. The log includes the episode number, reward, average reward over the last five episodes (Avg5), and epsilon value.

```
(venv) PS C:\Users\habha\Desktop\Reinforcement Learning\Assignment 3\RL_assignment3> python run_experiments.py
Ep 140 | Reward: -20.0 | Avg5: -19.60 | Eps: 0.05
Ep 150 | Reward: -20.0 | Avg5: -20.00 | Eps: 0.05
Ep 160 | Reward: -21.0 | Avg5: -20.60 | Eps: 0.05
Ep 170 | Reward: -19.0 | Avg5: -20.40 | Eps: 0.05
Ep 180 | Reward: -18.0 | Avg5: -18.60 | Eps: 0.05
Ep 190 | Reward: -21.0 | Avg5: -18.40 | Eps: 0.05
Ep 200 | Reward: -19.0 | Avg5: -18.00 | Eps: 0.05
Ep 210 | Reward: -16.0 | Avg5: -17.40 | Eps: 0.05
Ep 220 | Reward: -15.0 | Avg5: -15.00 | Eps: 0.05
Ep 230 | Reward: -18.0 | Avg5: -15.60 | Eps: 0.05
Ep 240 | Reward: -19.0 | Avg5: -16.00 | Eps: 0.05
Ep 250 | Reward: -15.0 | Avg5: -15.40 | Eps: 0.05
Ep 260 | Reward: -17.0 | Avg5: -17.60 | Eps: 0.05
```